

A rank collapse in a sedenion Module-LWE proposal, and why a uniform matrix does not repair it

Matthew Gifford

Independent researcher, Hollister, California, USA

October 2026

Abstract

We report a negative result on SLWE, a Module-LWE variant over the sixteen-dimensional sedenion algebra that the author proposed in an earlier paper with a claim of NIST Category 5 security. Its public matrix is built from orbits of an order-7 coordinate permutation (a “Singer cycle”). First, the flattened public matrix factors through a fixed 256×112 integer matrix of rank 76, so its rank over $\mathbb{F}_q$ is at most 76 for every module rank k , and exactly 76 for generic seeds once $k \geq 7$, whereas a uniform matrix of the same shape has rank $16k$. This invalidates the security claim at both published parameter sets. Recovering the noise then reduces to an LWE instance of dimension 76, which the lattice estimator places at about 2^{39} operations at the first set ($n = 512$, $q \approx 2^{32}$). Second, replacing the matrix by a uniform one does not help at that set (centred binomial noise with $\eta = 2$, ternary secrets of weight 64): the estimator gives about 2^{51} . Third, decryption at those parameters is deterministic: the noise is bounded by 258, far below $q/4$; a failure rate of $2^{-6.4 \times 10^{15}}$ reported earlier came from a Gaussian tail evaluated outside the support of a bounded sum. All computations are reproducible from public scripts.

1 Introduction

SLWE was proposed inside a larger, physics-motivated programme that uses the sedenions $\mathbb{S}$ (the Cayley–Dickson double of the octonions) and their zero divisors. The idea was to keep the shape of Module-LWE [10, 7] while replacing the ring by $\mathbb{S}_q = \mathbb{S} \otimes \mathbb{F}_q$ and to give the public matrix a symmetry inherited from the Fano plane: each row is the orbit of a random seed under a cyclic permutation σ of order 7. Correctness was obtained with a bilinear “norm” pairing rather than the (non-associative) algebra product.

Both versions of the paper that proposed SLWE [5] claimed NIST Category 5 security. Version 1 used $n = 512$ and $q \approx 2^{32}$. Version 2 found that modulus far too large (its primal estimate put version 1 below Category 1), moved to $k = 64$ and $q = 911$ ($n = 1024$), and claimed Category 5 from a primal estimate, leaving the dual and hybrid attacks to be checked. It kept the Singer-orbit matrix, whose rank it had tested only at $k = 4$. The project’s working specification [6] recorded on May 10, 2026 that this matrix is rank-deficient at $k = 32$, and replaced it by a uniform one.

This note shows that neither version is secure and records three facts that may be useful beyond this proposal.

1. A coordinate permutation applied to seeds produces a matrix that is linear in the seeds through a fixed “design matrix”. The rank of that design matrix bounds the rank of the public matrix for every module rank (Section 3).
2. Non-associativity of the algebra plays no role once decryption uses a bilinear form: every map the attacker needs is $\mathbb{F}_q$ -linear on $\mathbb{F}_q^{16k}$.

3. At $q \approx 2^{32}$ with noise of standard deviation about 1, LWE in dimension 512 is easy whatever the matrix (Section 4). The modulus-to-noise ratio, not the algebra, governs security. Version 2 of [5] reached the same conclusion with a primal estimate; Section 4 confirms it with the standard estimator.

2 The scheme

Let $\mathbb{S}_q$ be the sedenions over $\mathbb{F}_q$ with basis $e_0, \dots, e_{15}$, built by four Cayley–Dickson doublings $(a, b)(c, d) = (ac - \bar{d}b, da + b\bar{c})$. Write $L(a) \in \mathbb{F}_q^{16 \times 16}$ for left multiplication by a and $\langle u, v \rangle = \text{Re}(\bar{u}v) = \sum_t u_t v_t$ for the norm pairing, extended to $\mathbb{S}_q^k$ by summing over components. The identity $L(a)^\top = L(\bar{a})$ gives $\langle As, r \rangle = \langle s, A^H r \rangle$ with $A_{ij}^H = \overline{A_{ji}}$, which is what decryption needs.

- **KeyGen.** $s \in \mathbb{S}_q^k$ is sparse ternary (weight h_s). $A \in \mathbb{S}_q^{k \times k}$ has $A_{ij} = \sigma^j(\text{seed}_i)$ with independent uniform seeds, where σ moves the coefficient of e_t to $e_{(t \bmod 7)+1}$ for $t = 1..7$, moves that of e_{t+8} to $e_{(t \bmod 7)+9}$, and fixes e_0, e_8 . The noise e is centred binomial (or confined to the four-dimensional kernel of a zero divisor [9]). The public key is $(A, b = As + e)$.
- **Encaps.** r is sparse ternary (weight h_r), with $c_1 = A^H r + e_1$ and $c_2 = \langle b, r \rangle + e_2 + m \lfloor q/2 \rfloor$.
- **Decaps.** Round $c_2 - \langle s, c_1 \rangle = m \lfloor q/2 \rfloor + N$ with $N = \langle e, r \rangle - \langle s, e_1 \rangle + e_2$.

Version 1’s parameters are $k = 32$ (so $n = 16k = 512$), $q = 4294977961$ (the least prime $\equiv 1 \pmod{455}$ above 2^{32}), $h_s = h_r = 64$, and noise $\text{CBD}(\eta = 2)$. Version 2’s are $k = 64$ ($n = 1024$) and $q = 911$ (the least prime $\equiv 1 \pmod{455}$), with the same weights and noise.

3 The rank collapse

Let $M \in \mathbb{F}_q^{16k \times 16k}$ be the matrix of $s \mapsto As$, i.e. $M_{16i+u, 16j+c} = \text{coefficient } u \text{ of } \sigma^j(\text{seed}_i) e_c$.

Proposition 1. *Let $N \in \mathbb{Z}^{256 \times 112}$ have entries $N_{16m+u, 16j+c} = \text{coefficient } u \text{ of } \sigma^j(e_m) e_c$ ($0 \leq m, u, c \leq 15, 0 \leq j \leq 6$). Then $M = (S \otimes I_{16}) N_k$, where $S \in \mathbb{F}_q^{k \times 16}$ holds the seeds and N_k consists of the column blocks $j \bmod 7$ of N , $j = 0, \dots, k-1$. Hence $\text{rank } M \leq \text{rank } N$ for every k . Moreover $\text{rank}_{\mathbb{Q}} N = 76$, and $\text{rank } N = 76$ modulo 911 and modulo q .*

Proof. The product is bilinear and σ is linear, so $\sigma^j(\text{seed}_i) e_c = \sum_m (\text{seed}_i)_m \sigma^j(e_m) e_c$. That is the stated factorisation. Since $\sigma^7 = \text{id}$, column block j of N_k equals block $j \bmod 7$, so N_k has at most the columns of N . The ranks of N were computed exactly (Section 7). $\square$

The bound is attained. Table 1 lists exact ranks from two independent implementations (one built on the proposal’s reference code, one written blind from a textual specification), with three seed sets per case and a uniform control. The collapse starts at $k = 5$; from $k = 7$ on the rank is 76 for every seed set tried. Version 2’s parameters, $k = 64$ over $\mathbb{F}_{911}$, are in the table: the rank is 76, not 1024. Its own rank test used $k = 4$, where the matrix has 64 columns, below the cap of 76, so the collapse could not show there.

Remark 1. σ is not an automorphism of the sedenion product, so the “symmetry” of A is a relabelling of coordinates. The argument above uses only linearity.

Corollary 1. *Let V be the column space of M ($\dim V = \rho \leq 76$). Then $e \in b + V$. Choose ρ coordinates I on which a basis of V is invertible. Then $e_{\bar{I}} = -C e_I + (b_{\bar{I}} + C b_I)$ for a public C , an LWE instance with secret e_I of dimension ρ and $16k - \rho$ samples. The same holds for e_1 and the column space of M^H . Given e and e_1 , any r' with $M^H r' = c_1 - e_1$ gives $c_2 - \langle b - e, r' \rangle = m \lfloor q/2 \rfloor + \langle e, r \rangle + e_2$, which reveals m .*

k	1	2	3	4	5	6	7	8	16	32	64
$16k$	16	32	48	64	80	96	112	128	256	512	1024
rank M (Singer)	16	32	48	64	72	74	76	76	76	76	76
rank M (uniform)	16	32	48	64	80	96	112	128	256	512	1024

Table 1: Exact ranks over $\mathbb{F}_{911}$ (the values for $k \in \{8, 16, 32, 64\}$ are identical over $\mathbb{F}_q$ with $q = 4\,294\,977\,961$, for three seed sets each, in both implementations).

For $\rho = 76$, $q = 4\,294\,977\,961$, CBD(2) secret and noise and $m = 436$ samples, the lattice estimator [2, 3] reports $2^{38.6}$ operations (uSVP and BDD). Its block size is at the floor of its model ($\beta = 40$), so reduction at the scale of LLL already suffices. Distinguishing (A, b) from uniform is easier still: it needs only Gaussian elimination, because A itself is rank-deficient. At version 2’s parameters the same reduction gives dimension 76 with 948 samples; we did not run the estimator there.

4 A uniform matrix does not repair version 1’s parameters

The working specification [6] replaced the Singer matrix by a uniform one and kept version 1’s parameters. Table 2 shows the estimator’s default attack set at those parameters with a uniform matrix. The cheapest attack costs about 2^{51} . An independent implementation of the primal “2016 estimate” [1] (smallest β with $\sigma\sqrt{\beta} \leq \delta_\beta^{2\beta-d-1} \text{Vol}^{1/d}$, the secret rescaled to the noise width) gives $\beta = 72$, $d = 937$, a core-SVP cost of $2^{0.292\beta} = 2^{21}$ [4], the same as the estimator’s own rough uSVP figure. The cause is the modulus: with $q/\sigma \approx 2^{32}$, a block size near 70 already exposes the short vector in dimension about 940.

q	cheapest attack ($\log_2$ rop)	uSVP ($\log_2$ rop, β)
4 294 977 961 (version 1)	51.0 (BDD)	52.1, $\beta = 72$
2731	114.4 (dual hybrid)	122.5, $\beta = 330$
911	122.2 (BKW); 127.4 without BKW	136.9, $\beta = 383$

Table 2: Lattice estimator [3], commit 53da598, default attacks and cost model [8], at $n = 512$, $m = 512$, ternary secret of weight 64 and CBD(2) noise, with a uniform public matrix. The columns $q = 911$ and $q = 2731$ (the two smallest primes $\equiv 1 \pmod{455}$) are for orientation only; version 2 pairs $q = 911$ with $n = 1024$.

Even the smallest admissible prime falls short of 2^{128} at $n = 512$ with these distributions, so a secure instance would need a larger dimension or wider noise. We have not searched that space. Version 2’s parameters ($n = 1024$, $q = 911$) are not in the table: with the Singer matrix they fall to Section 3, and with a uniform matrix the scheme becomes generic Module-LWE with sparse secrets, which we did not estimate.

5 Correctness

Each coordinate of e , e_1 and e_2 lies in $[-2, 2]$. The confined zero-divisor noise has the same bound, since the kernel bases have entries in $\{0, \pm 1\}$ and the noise is $2 \sum_j \alpha_j k_j$ with $\alpha_j \in \{-1, 0, 1\}$. So

$$|N| \leq 2h_r + 2h_s + 2 = 258 < \lfloor q/4 \rfloor \approx 1.07 \times 10^9,$$

and decryption never fails at version 1’s parameters. It never fails for any $q > 1032$; at version 2’s $q = 911$ it can fail, though rarely. For CBD noise the distribution is in fact exact: given r and s , N is a sum of $2(h_r + h_s + 1)$ independent symmetric binomial draws,

i.e. $N \sim \text{CBD}(258)$. At $q = 911$ this gives $\Pr[|N| > 227] = 2^{-353.5}$ exactly, consistent with version 2's Gaussian bound of 2^{-295} .

An earlier analysis reported $\log_2 \text{DFR} \approx -6.4 \times 10^{15}$. That number is the Gaussian tail at $z = q/(4\sigma_N) \approx 9.5 \times 10^7$ with $\sigma_N \approx 11.4$. It evaluates the tail about 4×10^6 times beyond the largest value N can take. At that modulus the correct statement is $\text{DFR} = 0$.

6 Conclusion

The Singer-orbit public matrix makes SLWE insecure at every module rank, including both published parameter sets. A uniform matrix does not rescue version 1's parameters, which reach about 2^{51} . The scheme's correctness was never in question. The general lesson is elementary but easy to miss. A structured matrix built by acting on random seeds is the image of a fixed design matrix, and that matrix's rank should be the first thing checked.

7 Reproducibility

All scripts, outputs and checksums are in the public repository `gifgaf0/gifgaf0.github.io` at commit

`b09929ada97cc3c5bfc0f5f327518cd82e342fea`

in the directory `audit_followup/phase_c/C1/`:

- the two rank implementations: `leg1/`, and `leg2/`, written blind from the pre-registration `C1_PREREG.md`;
- their comparison and the computation of rank N : `c1_compare.py`;
- the estimator runs: `estimator/`, using `passagemath 10.8.12` and `lattice-estimator` commit `53da5982597709ba0fdf94ea37a84d822310fd84`;
- the core-SVP cross-check: `c1_coresvp_crosscheck.py`;
- the noise bounds and exact tails: `c1_checks.py`.

Exact ranks were computed with `python-flint`.

Acknowledgements. AI agents (Anthropic's Claude) drafted this note and wrote and ran all of its computations, including both independent rank implementations; the author directed the work, reviewed the text, and takes responsibility for it.

References

- [1] E. Alkim, L. Ducas, T. Pöppelmann, P. Schwabe. Post-quantum key exchange — a new hope. *25th USENIX Security Symposium*, 2016, 327–343.
- [2] M. R. Albrecht, R. Player, S. Scott. On the concrete hardness of Learning with Errors. *J. Math. Cryptol.* 9(3) (2015), 169–203.
- [3] M. R. Albrecht et al. The lattice estimator. <https://github.com/malb/lattice-estimator>, commit `53da598`.
- [4] A. Becker, L. Ducas, N. Gama, T. Laarhoven. New directions in nearest neighbor searching with applications to lattice sieving. *SODA 2016*, 10–24.

- [5] M. Gifford. Fluid Lattice Topology: bridging $\text{PSL}(2,7)$ vacuum geometry and non-associative post-quantum cryptography. Zenodo, 2026, versions 1 and 2. doi:10.5281/zenodo.20078312.
- [6] M. Gifford. Sedenion Module-LWE and Φ -modular prime project: master reference document, v2.1 (May 10, 2026). File `tools/SLWE_Prime_Master_v2.md` in the repository of Section 7.
- [7] A. Langlois, D. Stehlé. Worst-case to average-case reductions for module lattices. *Des. Codes Cryptogr.* 75 (2015), 565–599.
- [8] MATZOV. Report on the security of LWE: improved dual lattice attack. 2022. doi:10.5281/zenodo.6412487.
- [9] R. G. Moreno. The zero divisors of the Cayley–Dickson algebras over the real numbers. arXiv:q-alg/9710013, 1997.
- [10] O. Regev. On lattices, learning with errors, random linear codes, and cryptography. *J. ACM* 56(6) (2009), article 34.